

Lecture 1

Neural networks: flexible functions, gradient descent, and the generalization question

Deep Neural Networks · University of Warsaw · 2026/27

Course logistics

Announcements: Slack



All announcements are on Slack: scan the code to join the workspace, then join **#dnn-general**.

[invitation link](#)

Materials

mim-ml-teaching.github.io/public-dnn

Grading

The grading rules will be announced in a week.

Entry quiz (NumPy and other prerequisites)

- the questions are **public** (on the materials page): use them to check yourself
- quiz at the **5th labs** (Oct 30 to Nov 6): a random subset of the same questions, no notes
- 10 questions, **10 points** out of about 100 to 120 in the course (roughly half a grade)

ML Bootcamp (optional)

- **Mon & Wed, 18:15**, rooms 3043/3044, first four weeks; come to any subset
- refresh the prerequisites: NumPy, Pandas, gradients, linear regression
- or just come with questions, e.g. about previous labs

In USOS the bootcamp appears as groups 10 and 11; they are unrelated to the Tuesday 12:15 and 14:15 lab groups.

Today in one sentence

- a neural network is a **highly customizable function**
- with a **huge number of parameters**
- **gradient descent** iteratively adjusts these parameters to reduce the training loss
- the result **generalizes** far better than it has any right to

1 · the function

what a network computes, and how flexible it is

2 · fitting it

loss, gradient descent, SGD, gradients by hand


3 · generalization

why does it work on data it has never seen?

Part 1

Programs we cannot write

What is inside these boxes?



Question Could you write any of these programs by hand?

Each box is a function

$$f : \mathcal{X} \rightarrow \mathcal{Y}$$

- we cannot write f down...
- ...but we can collect **examples** (x_i, y_i)

The deep learning recipe

1. Choose a neural-network architecture that composes multiple learned layers with nonlinear activations.
2. Fit its parameters to data by minimizing a training loss, typically using gradient-based optimization.

Not a new idea

Linear regression is the same recipe with a tiny family:

$$\begin{aligned}\hat{y} = f_{\theta}(\mathbf{x}) &= \theta_0 + \sum_{j=1}^d \theta_j x_j \\ &= \theta_0 + \theta_1 x_1 + \cdots + \theta_d x_d.\end{aligned}$$


What changed: the family is enormous, and we still manage to find good θ .

A highly customizable function

Where it began: the perceptron (1957)


Frank Rosenblatt's model of a neuron:

$$y = \left[w_1 x_1 + \cdots + w_k x_k + b \geq 0 \right] \in \{0, 1\}$$



- a **weighted vote** of the inputs, then a yes/no decision
- the weights w_i and bias b are **knobs** — already a "function with knobs"
- a linear model followed by a threshold: the same formula as today's neuron, only the step function was later replaced


The Mark I Perceptron (Cornell, 1958–60)



Training, by hand:


1. put a printed letter in front of the camera
2. if the machine answers wrongly, the trainer flips a switch to **force** the correct answer
3. electric motors turn the potentiometers that store the weights

Result: one typeface learned without error; 85% correct once a second typeface was added.



Mark I Perceptron, Cornell Aeronautical Laboratory. US Navy photo, 1960, public domain.

What one perceptron can — and cannot — do



$$y = [1 x_1 + 1 x_2 - 0.5 \geq 0]$$

target: **AND** · correct: **2/4**

move the line until all four are correct

One perceptron draws **one line**: XOR is impossible (Minsky & Papert, Perceptrons, 1969). Two layers solve it — next slide: in fact, any Boolean function.

Two layers compute any Boolean function

Example: $f(x_1, x_2, x_3) = 1$ iff **exactly one** input is 1.


x_1	x_2	x_3	f
0	0	0	0
0	0	1	1
0	1	0	1
0	1	1	0
1	0	0	1
1	0	1	0
1	1	0	0
1	1	1	0

click a row

— weight +1

— weight -1

thick: input is 1



Recipe (OR of ANDs): one hidden perceptron per input pattern with $f = 1$ — weight +1 for inputs that must be 1, -1 for inputs that must be 0, threshold = number of 1s in the pattern, so it fires on that pattern only. The output perceptron is an OR: all weights 1, threshold 1. Catch: up to 2^k hidden units for k inputs (e.g. parity needs 2^{k-1}).

From perceptrons to the multi-layer perceptron

Two layers of perceptrons can compute any Boolean function (previous slide). But how do we **train** the hidden layer?

The Mark I left its hidden layer random: nobody knew how to train it.


Nudge any weight slightly: the step output does not change at all. Every partial derivative is **0** — gradient descent has nothing to follow.

Fix (1980s, backpropagation): replace the step by a **smooth** or piecewise-linear activation. Same weighted sum, same layers — now trainable by gradient descent.


The name stayed: a **multi-layer perceptron (MLP)** today means layers of such neurons — not of 1957 perceptrons.



step (perceptron)
derivative 0 everywhere
(undefined at 0)




sigmoid (1980s)
smooth: nudges have an effect



ReLU (2010s)
today's default

Why now?



Data

MNIST (1998): 60,000 images → ImageNet (2009): 1.2 million → today's language models: ~15 trillion tokens of text.

Compute

AlexNet: 2 gaming GPUs for about a week. Since 2010, the compute used to train frontier models has grown roughly **4–5× per year**.


Algorithms & software

ReLU, better initialization, normalization, residual connections, Adam, attention — and frameworks that compute gradients automatically (next lecture).

Neuron, layer, network

neuron


weighted sum + nonlinearity



$$h = \text{ReLU}(\mathbf{w}^\top \mathbf{x} + b)$$
$$f : \mathbb{R}^2 \rightarrow \mathbb{R}$$

layer

m neurons side by side




$$\mathbf{h} = \text{ReLU}(W\mathbf{x} + \mathbf{b})$$
$$f : \mathbb{R}^2 \rightarrow \mathbb{R}^m$$

$W \in \mathbb{R}^{m \times 2}$ has rows $\mathbf{w}_1^\top, \dots, \mathbf{w}_m^\top$

network

layers composed



$$\mathbf{h}_1 = \text{ReLU}(W_1\mathbf{x} + \mathbf{b}_1)$$


$$\mathbf{h}_2 = \text{ReLU}(W_2\mathbf{h}_1 + \mathbf{b}_2)$$

$$f_\theta(\mathbf{x}) = W_3\mathbf{h}_2 + \mathbf{b}_3$$

One neuron, one layer

neuron: $h = \text{ReLU}(\mathbf{w}^\top \mathbf{x} + b)$

layer, outputs summed: $\sum_k v_k \text{ReLU}(\mathbf{w}_k^\top \mathbf{x} + b_k)$



one neuron: $h = \text{ReLU}(\mathbf{w}^\top \mathbf{x} + b)$


a hinge: zero on one side of the purple line, a tilted plane on the other

parameters: **3**

Here the neuron is a function $f : \mathbb{R}^2 \rightarrow \mathbb{R}$, and so is the layer once its m outputs are summed with weights. **Input:** a point $\mathbf{x} = (x_1, x_2)$ on the floor plane. **Output:** one number, drawn as the height of the surface above that point.

Question Why do we need the nonlinearity at all? → turn the ReLU off: any number of units collapses to **one plane**.

One hidden layer = a piecewise-linear function




hidden units: **6**

parameters: **19**

linear pieces: ≤ 7

loss (MSE): **0.0059**

training steps: 3000



$$f(x) = c + \sum_{k=1}^n a_k \text{ReLU}(w_k x + b_k)$$

Unit k is **off** ($= 0$) on one side of its joint $x = -b_k/w_k$ and **linear** on the other.

Between two neighbouring joints every unit is either off or linear, so f is a sum of linear functions — **linear**.

At joint k the slope jumps by $a_k w_k$. So n units $\rightarrow$ at most n joints $\rightarrow$ at most $n + 1$ linear pieces. More units $\rightarrow$ more joints $\rightarrow$ any shape.

Depth: folding space



layers: **3**

ReLU units: **6**

linear pieces: **8**

one hidden layer needs ≈ 8
units for this function

Each layer folds $[0, 1]$ onto itself and so **doubles** the number of joints: k layers of 2 units give 2^k linear pieces. A single hidden layer needs a unit for every joint — about 2^k units for the same function.

Depth lets the same number of parameters encode **exponentially many joints** in the shape of the function: width adds joints, depth multiplies them.

Universal approximation

Theorem (Cybenko 1989; Hornik 1991). A network with one hidden layer can approximate any continuous function on a compact set arbitrarily well — given enough hidden units.


Reassuring, but it says much less than it seems:

question	does the theorem answer it?
can some θ represent the function?	yes
how many units are needed?	no — possibly astronomically many (depth helps: previous slide)
can we find that θ from data?	no
will it generalize to new inputs?	no

Polynomials and lookup tables are universal approximators too. Expressiveness is not the explanation.

How many knobs?

■ total parameters ■ used per token (mixture of experts)



Each parameter is one real number the training procedure has to set. The largest models are **mixtures of experts**: only a fraction of the parameters is used for each token. Frontier closed models (GPT, Claude, Gemini) do not publish their sizes.

Fitting it: (stochastic) gradient descent

Learning = minimizing a loss

Training set $\{(\mathbf{x}_i, y_i)\}_{i=1}^n$, a loss per example ℓ :

$$L(\theta) = \frac{1}{n} \sum_{i=1}^n \ell(f_{\theta}(\mathbf{x}_i), y_i)$$


- regression: $\ell = (\hat{y} - y)^2$
- classification: $\ell = -\log p_{\theta}(y \mid \mathbf{x})$ (cross-entropy)

A landscape in millions of dimensions

L is a function of all parameters. Training = walking downhill on it.

It is **not convex**: no guarantee of finding the global minimum. In practice that hardly matters — one of today's surprises.

Gradient descent vs stochastic gradient descent



- GD (all 64 points): loss **0.063**

- SGD (batch 4): loss **0.046**


steps: 60 · cost per step: GD 64 vs
SGD 4 gradient terms

GD: $\theta \leftarrow \theta - \eta \nabla L(\theta)$ — uses **all** n examples per step.

SGD: $\theta \leftarrow \theta - \eta \nabla \frac{1}{|B|} \sum_{i \in B} \ell_i(\theta)$ — a random **minibatch** B :
noisy, but $\frac{n}{|B|}$ times cheaper.

Watch a network learn

pick a scenario above, or set the parameters by hand



● class 0 ● class 1 · filled: training · ○ test
■ ■ background: network's prediction
(stronger = more confident)

parameters: 205 · epoch 300

accuracy: train 99% · test 98%



The same loop trains everything

```
initialize  $\theta$  randomly
repeat:
  B  $\leftarrow$  random minibatch
  L  $\leftarrow$  loss of  $f_{\theta}$  on B
  g  $\leftarrow \nabla_{\theta} L$  # how?
   $\theta \leftarrow \theta - \eta \cdot g$ 
```


From the spiral above to trillion-parameter language models: **the same loop**.

Everything hinges on one line: **computing the gradient** of the loss with respect to every parameter.


Let us do it once by hand.

Which way should each knob turn?

Question If I change one weight by a small ε , what happens to the loss?



forward pass shown for $\varepsilon = +0.5$



$$L(\theta) = 0.5$$

$$\varepsilon = +0.5$$

$$(L(w+\varepsilon) - L)/\varepsilon = 2.06250$$

$$\varepsilon = -0.5$$

$$(L - L(w-\varepsilon))/\varepsilon = 0.93750$$

From nudges to derivatives

The ratio we measured, in the limit, is the **partial derivative**:

$$\frac{\partial L}{\partial w} = \lim_{\varepsilon \rightarrow 0} \frac{L(\dots, w + \varepsilon, \dots) - L(\dots, w, \dots)}{\varepsilon}$$

change one parameter, keep all others fixed

The **gradient** collects all of them:

$$\nabla_{\theta} L = \left(\frac{\partial L}{\partial \theta_1}, \dots, \frac{\partial L}{\partial \theta_P} \right)$$

and for a small change Δ of all parameters at once: $L(\theta + \Delta) \approx L(\theta) + \nabla_{\theta} L^{\top} \Delta$

What it tells us


- **sign**: which way to turn each knob to reduce the loss
- **magnitude**: how sensitive the loss is to that knob (0: no effect right now)
- $-\nabla_{\theta} L$ is the direction of steepest descent $\rightarrow$ the gradient step $\theta \leftarrow \theta - \eta \nabla_{\theta} L$

Why not just nudge?

Nudging = one extra forward pass **per parameter**:
669,706 forward passes for one gradient of a small MNIST classifier. The **chain rule** gets all of them from one forward and one backward pass.

Computing the gradient: the chain rule

0/9 · The network and one training example



$$\mathbf{x} = (1, 2), y = 1 \quad \mathbf{z} = W_1 \mathbf{x} + \mathbf{b}_1, \mathbf{h} = \text{ReLU}(\mathbf{z}), \hat{y} = \mathbf{v}^\top \mathbf{h} + b_2, L = \frac{1}{2}(\hat{y} - y)^2$$

This does not scale by hand

- our tiny network: **9** parameters, a page of algebra
- a small MNIST digit classifier: **669,706** parameters
- Kimi K3 (2026): **2,800,000,000,000**

But the computation was **mechanical**: local derivatives, multiplied along the paths from the loss back to each parameter, reusing shared intermediate results.

Next lecture: teach the computer to do exactly this — automatically, for any network. One backward pass costs about as much as one forward pass, however many parameters there are.

The generalization question: why does it work?

The classical worry

A network with one hidden layer of n units can **memorize any** n labelled examples: give each training point its own unit.


Classical statistics: a model that can fit anything learns nothing. More parameters than data → overfitting.

Our MNIST network

- 669,706 parameters
- 5,000 training images
- ≈ 134 parameters per example

It should overfit badly. Let us train it.

Experiment: true labels vs random labels




Same network, same optimizer, same images. With random labels it still reaches **100% training accuracy** — and, of course, chance-level test accuracy.

What does this tell us?

- The network **can** memorize arbitrary labels — classical capacity arguments say it is "too big".
- Yet on real data, SGD finds a solution that **generalizes**: among the many functions that fit the training set, it picks a good one.
- So generalization depends on the **data**, the **architecture** and the **training procedure** — not just on the number of parameters.

Zhang et al. (2017), "Understanding deep learning requires rethinking generalization".

Experiment: more parameters, less overfitting?




Double descent: classical U-curve on the left; a peak where the model can just barely fit the data; then test error falls again — far below the classical optimum.

Data	MNIST handwritten digits (28×28 grey images, 10 classes): 1,000 training images drawn at random, the full 10,000-image test set.
Architecture	one hidden layer of p ReLU units with random, fixed input weights; only the output layer is fitted, by least squares on one-hot labels. p swept from 10 to ~20,000; averaged over 5 random seeds.
Code	<code>experiments/double_descent.py</code> → <code>assets/lecture01/data/double_descent.json</code> ; reproduce with <code>python experiments/double_descent.py</code> (a few minutes on a laptop CPU).

Deep double descent (Nakkiran et al., 2019)

OPTIONAL



Redrawn from Fig. 1 (left) of Nakkiran, Kaplun, Bansal, Yang, Barak, Sutskever, Deep Double Descent: Where Bigger Models and More Data Hurt, ICLR 2020 (arXiv:1912.02292); values digitized from the published figure.

Setup

- ResNet18 with width $k = 1 \dots 64$ ($k = 64$: the standard ResNet18)
- CIFAR-10 with **15% label noise**, Adam, 4,000 epochs — no early stopping

Findings

- **Model-wise**: test error falls, rises around the point where the network can just fit the training set, then falls again — the widest network is best.
- **Epoch-wise**: for a fixed large model, test error over training time shows the same dip–rise–dip.
- **Sample-wise**: near the peak, **more training data can hurt** (4.5× more samples gave worse test error).
- Seen for ResNets, CNNs and transformers; **strongest with label noise**, small without it.

A priori, this should not work

- the loss landscape is **non-convex** — yet SGD finds near-zero training loss, almost from any initialization
- there are **many** zero-loss solutions, including ones that memorize noise — yet SGD lands on one that generalizes
- more parameters often **help** generalization, against the classical bias–variance intuition

What we understand (partly)

- **implicit bias** of gradient descent: among all fits it prefers "simple" ones (e.g. minimum norm)
- **flat minima** and the noise of SGD
- **architecture** encodes assumptions about the data (convolutions, attention — later lectures)
- **scale**: bigger models + more data keep improving

No complete theory yet — deep learning is still largely an empirical science.

The rest of the course

SGD in its plain form is fragile. Most of the course is the engineering that makes it work reliably — and the architectures that make it work on real data.

next lecture	automatic differentiation: computing $\nabla_{\theta}L$ for any network (frameworks, PyTorch)
--------------	---

optimization	momentum, Adam, learning-rate schedules
--------------	---

regularization	weight decay, augmentation, dropout
----------------	-------------------------------------

trainability	initialization, normalization, residual connections
--------------	---

architectures	convolutional networks, attention and transformers, generative models
---------------	---

Summary

- Tasks we cannot program are still **functions**; we pick f_θ from a huge family and let data choose θ .
- A ReLU network is a **piecewise-linear** function; width adds pieces, depth multiplies them; universal approximation says representation is not the bottleneck.
- Training = minimizing a loss with **SGD**; it needs the gradient, computed by the **chain rule** (next lecture: automatically).
- Networks can **memorize random labels**, yet **generalize** on real data; more parameters can **reduce** test error. Why — only partly understood.

Feedback



Please spend a minute: your feedback helps me improve the next lecture.

tinyurl.com/mimuw-dnn-feedback

Anonymous, no login, three questions.